

# Titanic Dataset Analysis

November 13, 2017

## 1 Predict survival of the Titanic

### 1.0.1 Fundamental analysis on the famous Titanic Dataset using the logistic Regression and Visualization with seaborn.

Importing the train and test dataset

```
In [1]: import numpy as np
import pandas as pd
import csv

train = pd.read_csv('G:/My Data Science files/Logistic Regression/train.csv')
test = pd.read_csv('G:/My Data Science files/Logistic Regression/test.csv')
```

```
In [2]: train.head()
```

```
Out[2]:
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

```
In [3]: train.describe()
```

```
Out [3]:
```

	PassengerId	Survived	Pclass	Age	SibSp	\
count	891.000000	891.000000	891.000000	714.000000	891.000000	
mean	446.000000	0.383838	2.308642	29.699118	0.523008	
std	257.353842	0.486592	0.836071	14.526497	1.102743	
min	1.000000	0.000000	1.000000	0.420000	0.000000	
25%	223.500000	0.000000	2.000000	20.125000	0.000000	
50%	446.000000	0.000000	3.000000	28.000000	0.000000	
75%	668.500000	1.000000	3.000000	38.000000	1.000000	
max	891.000000	1.000000	3.000000	80.000000	8.000000	

	Parch	Fare
count	891.000000	891.000000
mean	0.381594	32.204208
std	0.806057	49.693429
min	0.000000	0.000000
25%	0.000000	7.910400
50%	0.000000	14.454200
75%	0.000000	31.000000
max	6.000000	512.329200

```
In [4]: # To find the number of missing values in a dataframe
```

```
train.isnull().sum()
```

```
Out [4]: PassengerId      0
Survived      0
Pclass        0
Name          0
Sex           0
Age          177
SibSp         0
Parch         0
Ticket        0
Fare          0
Cabin        687
Embarked      2
dtype: int64
```

```
In [5]: # dropping the irrelevant columns
```

```
train1 = train.drop(['PassengerId', 'Name', 'Ticket', 'Cabin'], axis=1)
```

```
# nice way to encode the categorical values
```

```
cleanup_nums = {"Pclass": {1: "High Class", 2: "Medium Class", 3: "Lower Class"}}
train1.replace(cleanup_nums, inplace=True)
train1.head()
```

```
Out [5]:
```

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	0	Lower Class	male	22.0	1	0	7.2500	S

1	1	High Class	female	38.0	1	0	71.2833	C
2	1	Lower Class	female	26.0	0	0	7.9250	S
3	1	High Class	female	35.0	1	0	53.1000	S
4	0	Lower Class	male	35.0	0	0	8.0500	S

In [6]: *# grouping the age and fare*

```
import matplotlib.pyplot as plt
import seaborn as sns
```


```
%matplotlib inline
```

```
age = train1['Age'].dropna(axis=0)
age.to_string
```

```
sns.distplot(age,bins=10,kde=True,rug=True)
```

```
# there are two ways to handle this missing values either we can impute the data using
# we can remove those values
```


Out [6]: <matplotlib.axes.\_subplots.AxesSubplot at 0x28e311ec0f0>



In [7]: *# distribution of the fare*


```
fare = train1['Fare']
fare.to_string
sns.distplot(fare,kde=True,rug=True)
```

Out[7]: <matplotlib.axes.\_subplots.AxesSubplot at 0x28e31511d0>





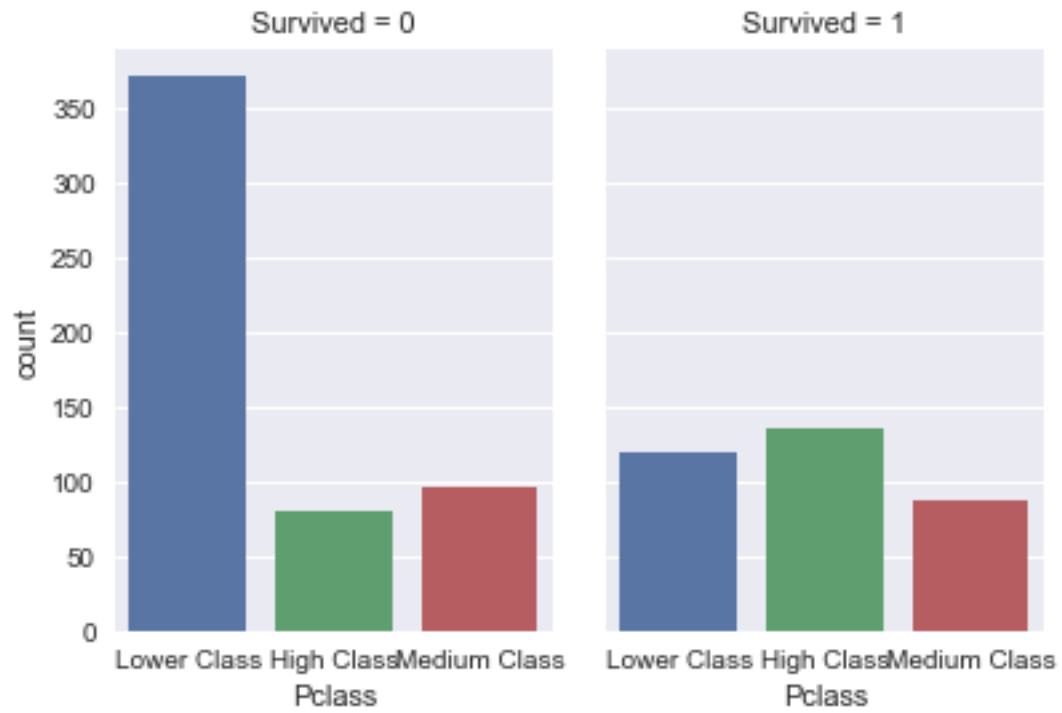
In [8]: sns.boxplot(fare)


Out[8]: <matplotlib.axes.\_subplots.AxesSubplot at 0x28e32695eb8>



```
In [9]: sns.set(style="darkgrid")
bar = sns.countplot(x="Pclass", data=train1)
g = sns.factorplot(x="Pclass", col="Survived", data=train1, kind="count", size=4, aspect=
h = sns.factorplot(x="Sex", col="Survived", data=train1, kind="count", size=4, aspect=
embarked = sns.factorplot(x="Embarked", col="Survived", data=train1, kind="count", size=
```







In [10]: *# importing the age through regression analysis*

```
#target = train1.loc[:, 'Survived']
#data_wo_survived = train1.iloc[:, 1:]
```

```
agetest = train1[train1['Age'].isnull()]
```

```
train1['Age'] = train1['Age'].fillna(9999)
```

```
agetrain = train1[train1['Age'] != 9999]
```

```
agetrain.describe()
```

```
# running regression analysis
```

```
agetrain = pd.get_dummies(agetrain)
```

```
agetrain.columns
```

Out[10]: Index(['Survived', 'Age', 'SibSp', 'Parch', 'Fare', 'Pclass\_High Class',  
'Pclass\_Lower Class', 'Pclass\_Medium Class', 'Sex\_female', 'Sex\_male',

```

        'Embarked_C', 'Embarked_Q', 'Embarked_S'],
        dtype='object')

In [11]: agetest.columns

Out[11]: Index(['Survived', 'Pclass', 'Sex', 'Age', 'SibSp', 'Parch', 'Fare',
               'Embarked'],
               dtype='object')

In [12]: agetest = agetest.iloc[:, [0,1,2,4,5,6,7]]

In [13]: agetest.head()
         agetest = pd.get_dummies(agetest)

In [14]: agetest.columns

Out[14]: Index(['Survived', 'SibSp', 'Parch', 'Fare', 'Pclass_High Class',
               'Pclass_Lower Class', 'Pclass_Medium Class', 'Sex_female', 'Sex_male',
               'Embarked_C', 'Embarked_Q', 'Embarked_S'],
               dtype='object')

In [15]: from sklearn import linear_model
         from sklearn.metrics import mean_squared_error, r2_score
         survive_data_agetrain = agetrain.iloc[:,0]
         survive_data_agetest = agetest.iloc[:,0]

         agetest_wo_survive = agetest.iloc[:,1:]

         agetrain_Y = agetrain.iloc[:,1]
         agetrain_X = agetrain.iloc[:,2:]

         regr = linear_model.LinearRegression()
         regr.fit(agetrain_X, agetrain_Y)

         # The coefficients
         print('Coefficients: ', regr.coef_)
         print('Intercept: ', regr.intercept_)

         agetrain_Y_predict = regr.predict(agetrain_X)

         # Explained variance score: 1 is perfect prediction
         print('Variance score: %.2f' % r2_score(agetrain_Y, agetrain_Y_predict))

         # The mean squared error
         print("Mean squared error: %.2f"
               % mean_squared_error(agetrain_Y, agetrain_Y_predict))

```



```
# Plot outputs
```


```
plt.scatter(agetrain_Y, agetrain_Y_predict, color='blue')  
plt.xlabel('Actual Age')  
plt.ylabel('Predict Age')  
  
plt.show()
```

```
Coefficients:  [-3.89188815  -0.7039223  -0.0221094   8.75649705  -6.7568533  
               -1.99964375  -1.52891861   1.52891861 -12.99425641  -7.76204648  
               -10.25152765]
```

```
Intercept:  44.5411738167
```

```
Variance score: 0.25
```

```
Mean squared error: 157.97
```



```
In [16]: # Using Statsmodels library to find the linear regression using OLS method.
```

```
import statsmodels.api as sm  
import matplotlib.pyplot as plt
```

```
model = sm.OLS(agetrain_Y, agetrain_X)  
results = model.fit()  
print(results.summary())
```

D:\Anaconda3\lib\site-packages\statsmodels\compat\pandas.py:56: FutureWarning: The pandas.core  
from pandas.core import datetools

```

                                OLS Regression Results
=====
Dep. Variable:                  Age      R-squared:                  0.250
Model:                        OLS      Adj. R-squared:              0.241
Method:                      Least Squares      F-statistic:              26.12
Date:                        Mon, 13 Nov 2017      Prob (F-statistic):        5.18e-39
Time:                        02:48:53      Log-Likelihood:            -2820.4
No. Observations:              714      AIC:                        5661.
Df Residuals:                  704      BIC:                        5707.
Df Model:                      9
Covariance Type:              nonrobust
=====
                                coef      std err          t      P>|t|      [0.025      0.975]
-----
SibSp                -3.8919      0.559      -6.960      0.000      -4.990      -2.794
Parch                -0.7039      0.631      -1.116      0.265      -1.943      0.535
Fare                 -0.0221      0.012      -1.861      0.063      -0.045      0.001
Pclass_High Class    26.5730      3.682       7.218      0.000      19.345      33.801
Pclass_Lower Class   11.0596      3.680       3.005      0.003       3.834      18.286
Pclass_Medium Class  15.8168      3.725       4.246      0.000       8.504      23.130
Sex_female           25.1958      5.402       4.664      0.000      14.589      35.803
Sex_male             28.2536      5.454       5.180      0.000      17.546      38.961
Embarked_C          -12.9943      9.056      -1.435      0.152     -30.774       4.785
Embarked_Q           -7.7620      9.352      -0.830      0.407     -26.122      10.598
Embarked_S          -10.2515      9.045      -1.133      0.257     -28.010       7.507
=====
Omnibus:                17.357      Durbin-Watson:              1.898
Prob(Omnibus):           0.000      Jarque-Bera (JB):           18.019
Skew:                   0.365      Prob(JB):                   0.000122
Kurtosis:               3.269      Cond. No.:                   4.81e+17
=====

```

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 1.23e-29. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

**1.0.2 Since the R2 is pretty low for the model, hence we will try for other regression methods.**

In [17]: `from sklearn.preprocessing import PolynomialFeatures`

```

x = [0]
y = [0]

```

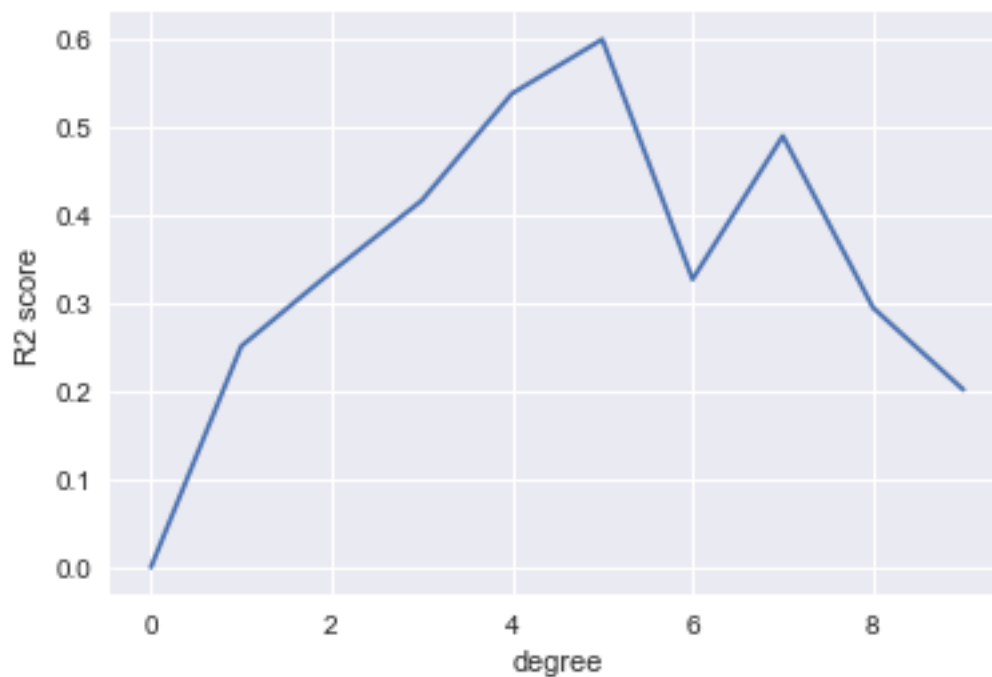
```

for i in np.arange(1,10,1):
    poly = PolynomialFeatures(degree=i,include_bias = False)
    X_poly = poly.fit_transform(agetrain_X)

    lin_reg = linear_model.LinearRegression()
    lin_reg.fit(X_poly,agetrain_Y)
    agetrain_Y_pred = lin_reg.predict(X_poly)
    x.append(i)
    a = r2_score(agetrain_Y, agetrain_Y_pred)
    y.append(a)

plt.plot(x,y)
plt.xlabel("degree")
plt.ylabel("R2 score")
plt.show()

```



**1.0.3 Since at polynomial degree = 5 we are getting a suitable R2 score**

**1.0.4 ,therefore we will be using polynimial regression of degree 5, to calculate the missing age.**

```

In [18]: poly = PolynomialFeatures(degree=5,include_bias = False)
         X_poly = poly.fit_transform(agetrain_X)

```

```
agetest_poly = poly.fit_transform(agetest_wo_survive)
```

```
lin_reg = linear_model.LinearRegression()
```

```
lin_reg.fit(X_poly,agetrain_Y)
```

```
agetest_Y_pred = lin_reg.predict(agetest_poly)
```

```
int_agetest_Y=[]
```

```
for i in agetest_Y_pred:
```

```
    int_agetest_Y.append(int(i))
```

```
In [19]: age_pred = np.array(int_agetest_Y)
```

```
age_predict = pd.DataFrame(age_pred,columns=['Age'])
```

```
age_predict = age_predict.reset_index()
```

```
agetest_wo_survive = agetest_wo_survive.reset_index()
```

```
survive_data_agetest = survive_data_agetest.reset_index()
```

```
agetest1 = pd.concat([survive_data_agetest,agetest_wo_survive,age_predict],axis=1)
```

```
In [20]: agetest1 = agetest1.drop(['index','level_0'],axis=1)
```

```
In [21]: agetest1.head()
```

```
Out [21]:
```

	Survived	SibSp	Parch	Fare	Pclass_High	Class	Pclass_Lower	Class	\
0	0	0	0	8.4583		0		1	
1	1	0	0	13.0000		0		0	
2	1	0	0	7.2250		0		1	
3	0	0	0	7.2250		0		1	
4	1	0	0	7.8792		0		1	

	Pclass_Medium	Class	Sex_female	Sex_male	Embarked_C	Embarked_Q	\
0		0	0	1	0	1	
1		1	0	1	0	0	
2		0	1	0	1	0	
3		0	0	1	1	0	
4		0	1	0	0	1	

	Embarked_S	Age
0	0	95
1	1	33
2	0	14
3	0	29
4	0	19

```
In [22]: # dropping the rows with negative age
```

```
agetest1 = agetest1.drop(agetest1[agetest1.Age < 0].index)
```

```
# dropping the rows with age > 100
```

```
agetest1 = agetest1.drop(agetest1[agetest1.Age > 100].index)
```

```
In [23]: # joining the dataframe with age predicted and dataframe with available age.
train_final = pd.concat([agetest1, agetrain])
```

```
In [24]: train_final.head()
#train_final.shape
```

```
Out [24]:
```

	Age	Embarked_C	Embarked_Q	Embarked_S	Fare	Parch	\
0	95.0	0	1	0	8.4583	0	
1	33.0	0	0	1	13.0000	0	
2	14.0	1	0	0	7.2250	0	
3	29.0	1	0	0	7.2250	0	
4	19.0	0	1	0	7.8792	0	

	Pclass_High Class	Pclass_Lower Class	Pclass_Medium Class	Sex_female	\
0	0	1	0	0	
1	0	0	1	0	
2	0	1	0	1	
3	0	1	0	0	
4	0	1	0	1	


  

	Sex_male	SibSp	Survived
0	1	0	0
1	1	0	1
2	0	0	1
3	1	0	0
4	0	0	1

### 1.0.5 Checking for collinearity

```
In [25]: import seaborn as sb
sb.heatmap(train_final.corr())
```

```
Out [25]: <matplotlib.axes._subplots.AxesSubplot at 0x28e35fdb630>
```



```
In [26]: train_final.corr()
```

```
Out[26]:
```

	Age	Embarked_C	Embarked_Q	Embarked_S	Fare	\
Age	1.000000	0.029511	-0.012771	-0.026256	0.102511	
Embarked_C	0.029511	1.000000	-0.138387	-0.803142	0.273551	
Embarked_Q	-0.012771	-0.138387	1.000000	-0.468976	-0.111601	
Embarked_S	-0.026256	-0.803142	-0.468976	1.000000	-0.181780	
Fare	0.102511	0.273551	-0.111601	-0.181780	1.000000	
Parch	-0.189084	-0.010164	-0.070661	0.053395	0.219426	
Pclass_High Class	0.365630	0.296553	-0.143063	-0.187558	0.590834	
Pclass_Lower Class	-0.321185	-0.156641	0.215463	0.016988	-0.424403	
Pclass_Medium Class	0.004650	-0.123952	-0.112235	0.179331	-0.108830	
Sex_female	-0.128774	0.084137	0.070473	-0.123688	0.191660	
Sex_male	0.128774	-0.084137	-0.070473	0.123688	-0.191660	
SibSp	-0.305680	-0.043817	-0.041573	0.066288	0.158250	
Survived	-0.081162	0.166877	-0.007497	-0.150330	0.274081	

	Parch	Pclass_High Class	Pclass_Lower Class	\
Age	-0.189084	0.365630	-0.321185	
Embarked_C	-0.010164	0.296553	-0.156641	
Embarked_Q	-0.070661	-0.143063	0.215463	

Embarked_S	0.053395	-0.187558	0.016988
Fare	0.219426	0.590834	-0.424403
Parch	1.000000	-0.006362	-0.007558
Pclass_High Class	-0.006362	1.000000	-0.629137
Pclass_Lower Class	-0.007558	-0.629137	1.000000
Pclass_Medium Class	0.016086	-0.293857	-0.558100
Sex_female	0.245051	0.103927	-0.153972
Sex_male	-0.245051	-0.103927	0.153972
SibSp	0.401943	-0.026003	0.041219
Survived	0.099255	0.290519	-0.335726


	Pclass_Medium Class	Sex_female	Sex_male	SibSp \
Age	0.004650	-0.128774	0.128774	-0.305680
Embarked_C	-0.123952	0.084137	-0.084137	-0.043817
Embarked_Q	-0.112235	0.070473	-0.070473	-0.041573
Embarked_S	0.179331	-0.123688	0.123688	0.066288
Fare	-0.108830	0.191660	-0.191660	0.158250
Parch	0.016086	0.245051	-0.245051	0.401943
Pclass_High Class	-0.293857	0.103927	-0.103927	-0.026003
Pclass_Lower Class	-0.558100	-0.153972	0.153972	0.041219
Pclass_Medium Class	1.000000	0.078398	-0.078398	-0.022929
Sex_female	0.078398	1.000000	-1.000000	0.131345
Sex_male	-0.078398	-1.000000	1.000000	-0.131345
SibSp	-0.022929	0.131345	-0.131345	1.000000
Survived	0.102714	0.540922	-0.540922	-0.000101

	Survived
Age	-0.081162
Embarked_C	0.166877
Embarked_Q	-0.007497
Embarked_S	-0.150330
Fare	0.274081
Parch	0.099255
Pclass_High Class	0.290519
Pclass_Lower Class	-0.335726
Pclass_Medium Class	0.102714
Sex_female	0.540922
Sex_male	-0.540922
SibSp	-0.000101
Survived	1.000000

```
In [27]: train_final = train_final.drop(['Embarked_S', 'Sex_female', 'Pclass_Lower Class'], axis=)
```

```
In [28]: sb.heatmap(train_final.corr())
```

```
Out[28]: <matplotlib.axes._subplots.AxesSubplot at 0x28e00167b38>
```



## 1.1 Logistic Regression

```
In [29]: from sklearn.linear_model import LogisticRegression
         from sklearn.metrics import accuracy_score
         X_target = train_final.iloc[:, -1]
         X_var = train_final.iloc[:, :-1]

         logistic = LogisticRegression()
         logistic.fit(X_var, X_target)
         X_pred = logistic.predict(X_var)

         print("Accuracy : %.2f" % accuracy_score(X_target, X_pred))
```

Accuracy : 0.81

```
In [30]: import statsmodels.api as sm

         logit = sm.Logit(X_target, X_var)

         result = logit.fit()
```



```
print (result.summary())
```

Optimization terminated successfully.

Current function value: 0.457173

Iterations 6

#### Logit Regression Results

Dep. Variable:	Survived	No. Observations:	856
Model:	Logit	Df Residuals:	847
Method:	MLE	Df Model:	8
Date:	Mon, 13 Nov 2017	Pseudo R-squ.:	0.3154
Time:	02:52:01	Log-Likelihood:	-391.34
converged:	True	LL-Null:	-571.62
		LLR p-value:	5.053e-73

	coef	std err	z	P> z	[0.025	0.975]
Age	-0.0139	0.005	-2.584	0.010	-0.024	-0.003
Embarked_C	0.6249	0.227	2.752	0.006	0.180	1.070
Embarked_Q	0.7134	0.328	2.173	0.030	0.070	1.357
Fare	0.0039	0.003	1.384	0.166	-0.002	0.009
Parch	0.0606	0.113	0.536	0.592	-0.161	0.282
Pclass_High Class	2.0512	0.307	6.681	0.000	1.449	2.653
Pclass_Medium Class	1.5197	0.229	6.638	0.000	1.071	1.968
Sex_male	-2.1921	0.183	-11.958	0.000	-2.551	-1.833
SibSp	-0.1463	0.109	-1.345	0.179	-0.360	0.067

```
In [31]: X_var = X_var.drop(['Fare', 'Parch', 'SibSp'], axis=1)
```

```
In [32]: import statsmodels.api as sm
```

```
logit = sm.Logit(X_target, X_var)
```

```
result = logit.fit()
```

```
print (result.summary())
```

Optimization terminated successfully.

Current function value: 0.459439

Iterations 6

#### Logit Regression Results

Dep. Variable:	Survived	No. Observations:	856
Model:	Logit	Df Residuals:	850
Method:	MLE	Df Model:	5
Date:	Mon, 13 Nov 2017	Pseudo R-squ.:	0.3120

Time: 02:52:21 Log-Likelihood: -393.28  
 converged: True LL-Null: -571.62  
 LLR p-value: 6.396e-75

	coef	std err	z	P> z	[0.025	0.975]
Age	-0.0137	0.005	-2.609	0.009	-0.024	-0.003
Embarked_C	0.6871	0.223	3.080	0.002	0.250	1.124
Embarked_Q	0.7290	0.326	2.236	0.025	0.090	1.368
Pclass_High Class	2.2868	0.253	9.024	0.000	1.790	2.783
Pclass_Medium Class	1.5722	0.225	6.983	0.000	1.131	2.013
Sex_male	-2.2241	0.181	-12.262	0.000	-2.580	-1.869

### 1.1.1 odds ratio

```
In [36]: import numpy as np
         print (np.exp(result.params))
```

```
Age          0.986430
Embarked_C   1.987977
Embarked_Q   2.073054
Pclass_High Class  9.843000
Pclass_Medium Class 4.817069
Sex_male     0.108160
dtype: float64
```

### 1.1.2 odds ratio and 95%CI

```
In [37]: # odds ratios and 95% CI
         params = result.params
         conf = result.conf_int()
         conf['OR'] = params
         conf.columns = ['2.5%', '97.5%', 'OR']
         print (np.exp(conf))
```

	2.5%	97.5%	OR
Age	0.976358	0.996606	0.986430
Embarked_C	1.283922	3.078107	1.987977
Embarked_Q	1.094328	3.927116	2.073054
Pclass_High Class	5.989871	16.174747	9.843000
Pclass_Medium Class	3.098452	7.488952	4.817069
Sex_male	0.075800	0.154336	0.108160

```
In [38]: logistic = LogisticRegression()
         logistic.fit(X_var,X_target)
```

```
X_pred = logistic.predict(X_var)

print("Accuracy : %.2f" % accuracy_score(X_target, X_pred))
```

Accuracy : 0.79

## 1.2 Confusion Matrix

```
In [39]: from sklearn.metrics import confusion_matrix

confusion_matrix = confusion_matrix(X_target, X_pred)

confusion_matrix

Out[39]: array([[447,  77],
               [103, 229]], dtype=int64)
```

## 1.3 Classification Report

```
In [40]: from sklearn import metrics
         from sklearn.metrics import classification_report

print(classification_report(X_target, X_pred))
```

	precision	recall	f1-score	support
0	0.81	0.85	0.83	524
1	0.75	0.69	0.72	332
avg / total	0.79	0.79	0.79	856

```
In [43]: test.head()
         # dropping the irrelevant columns
         test1 = test.drop(['PassengerId', 'Name', 'Ticket', 'Cabin'], axis=1)

         # nice way to encode the categorical values
         cleanup_nums = {"Pclass": {1: "High Class", 2: "Medium Class", 3: "Lower Class"}}
         test1.replace(cleanup_nums, inplace=True)
         test1.head()

         test1 = pd.get_dummies(test1)

In [44]: test1.head()
         test1 = test1.drop(['Embarked_S', 'Sex_female', 'Pclass_Lower Class'], axis=1)

In [45]: test1 = test1.drop(['Fare', 'Parch', 'SibSp'], axis=1)
```

## 1.4 Predicting the values of test data

```
In [46]: test1 = test1.dropna()
```

```
y_pred = logistic.predict(test1)
```

## 1.5 Predicted Values

```
In [47]: y_pred
```

```
Out[47]: array([0, 0, 0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1,
               1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1,
               1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 0, 0, 1, 1, 1, 1, 1, 0, 1, 0, 1,
               1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0,
               0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
               1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
               1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1,
               1, 1, 1, 0, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
               1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
               1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1,
               1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 1, 1, 1,
               1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
               1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1,
               1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1,
               1, 1, 1, 1, 1, 1, 0, 1, 1, 1], dtype=int64)
```

```
In [48]: y = pd.DataFrame(y_pred, columns = ['Survived'])
         test_pred = pd.concat([y, test], axis=1)
```